

CDC Academy Data Science Certificate Program

Basics of R Programming Participant's Guide

April 2021



R Programming



BeVera

Learning OBJECTIVES: during these sessions, Our Data Science trainers will:

- R Programming Basics
 - Rstudio Setup
 - Data Types
 - Data Structures
 - Packages
 - Importing Excel data files
- Explain the significance of exploratory data analysis and data science
- Demonstrate basic data science techniques using the R programming language
 - Apply basic tools plots, graphs, to carry out EDA

DATA PREPARATION

DATA CLEANING

INCONSISTENT DATATYPES
MISSPELLED ATTRIBUTES
MISSING AND DUPLICATE VALUES

TRANSFORMATION



EXPLORATORY DATA ANALYSIS



DEFINES AND REFINES
THE SELECTION OF FEATURE
VARIABLES THAT WILL BE USED
IN THE MODEL DEVELOPMENT

DATA MODELING

simplilearn

KNN



NAIVE BAYES

DECISION TREE

VISUALIZATION AND COMMUNICATION



WHAT IS DATA SCIENCE?

DATA ACQUISITION


- WEB SERVERS
- LOGS
- DATABASES
- APIS
- ONLINE REPOSITORIES



WHY?...WHY?...WHY?....



DEPLOYS AND





Module 1: R Programming Basics

<https://cran.r-project.org/doc/manuals/R-intro.pdf>

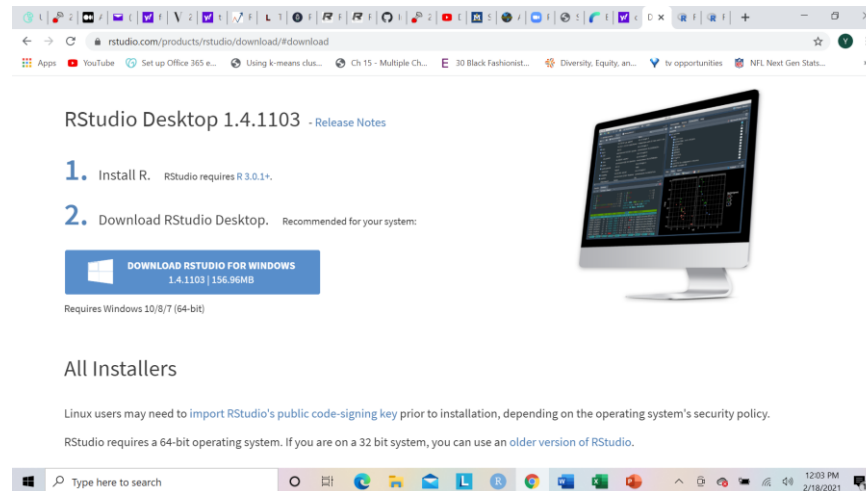
Setting up your machine

- Download a copy of [R](#) on your local computer from the Comprehensive R Archive Network (CRAN). You can choose between binaries for Linux, Mac and Windows.
- Install one of R's integrated development environment (IDE), [RStudio](#), which makes R coding much easier and faster as it allows you to type multiple lines of code, handle plots, install and maintain packages and navigate your programming environment much more productively.
- Install all suggested packages or dependencies including GUI.

<https://cran.r-project.org/doc/manuals/R-intro.pdf>

Setting up your machine

- <https://www.r-project.org/>
- <https://rstudio.com/products/rstudio/download/#download>
- <https://cran.r-project.org/mirrors.html>



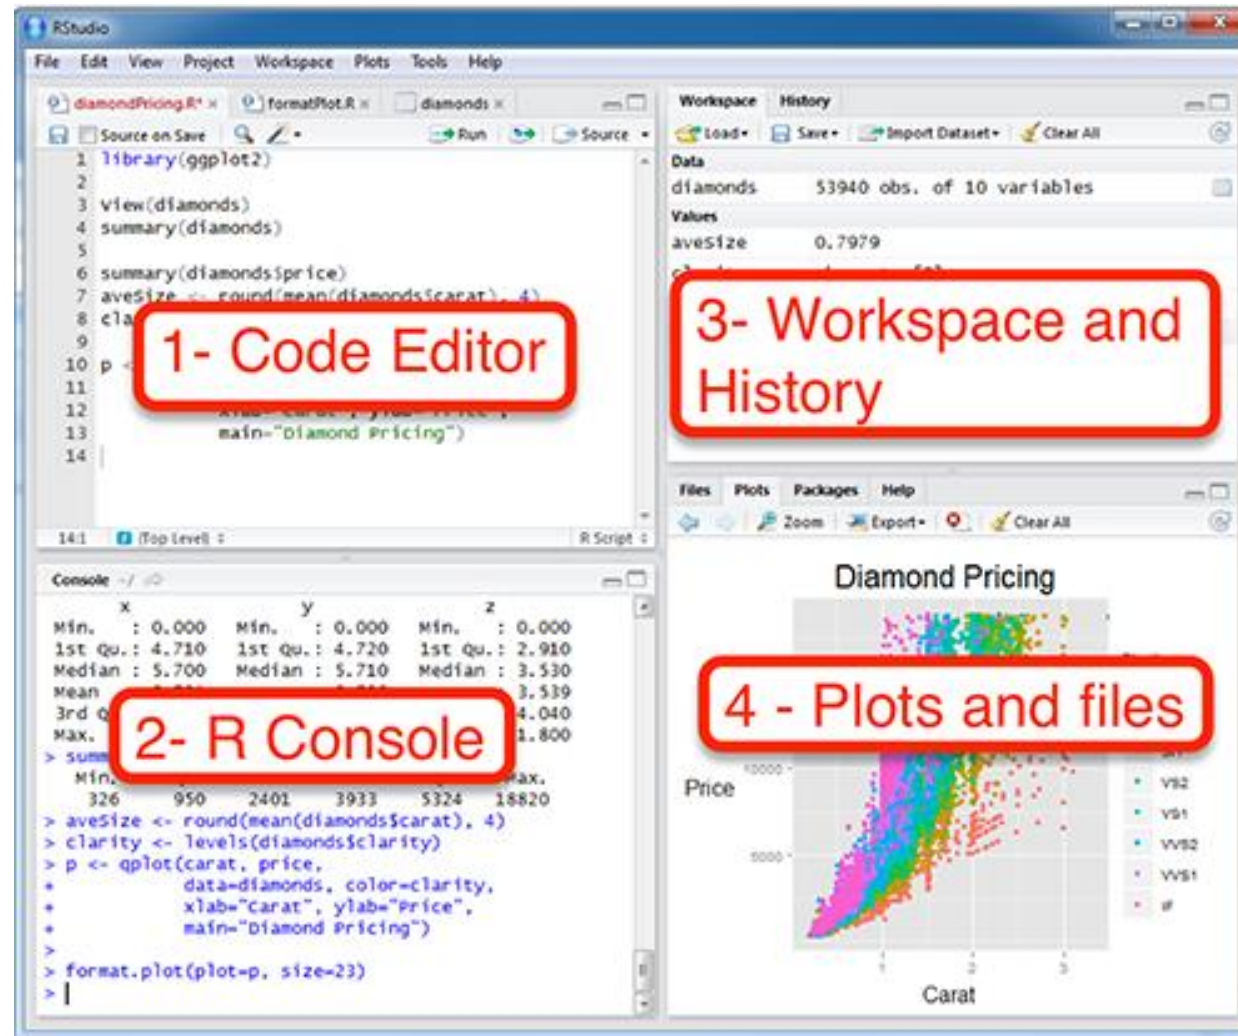
RStudio is a four pane work-space for 1) creating file containing R script, 2) typing R commands, 3) viewing command histories, 4) viewing plots and more.

1.Top-left panel:

- Code editor allowing you to create and open a file containing R script.
- The R script is where you keep a record of your work. R script can be created as follow: File → New → R Script

2.Bottom-left panel:

- R console for typing R commands



3.Top-right panel:

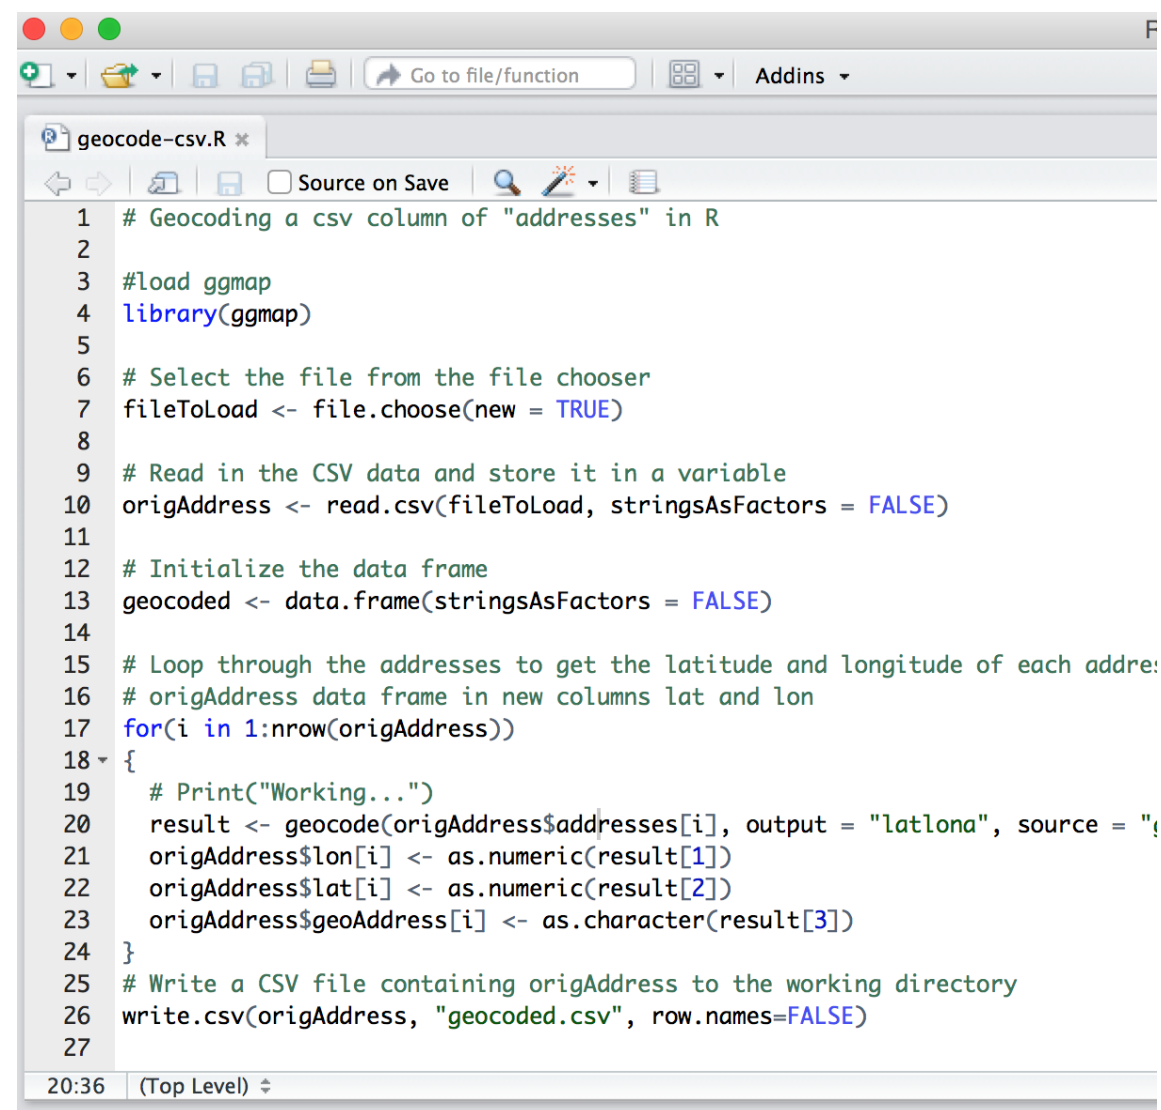
- Workspace tab: shows the list of R objects you created during your R session
- History tab: shows the history of all previous commands

4.Bottom-right panel:

- Files tab: show files in your working directory
- Plots tab: show the history of plots you created. From this tab, you can export a plot to a PDF or an image files
- Packages tab: show external R packages available on your system. If checked, the package is loaded in R.

R Script Files

Often you will do your programming by writing your programs in script files and then you execute those scripts at your command prompt with the help of R interpreter called Rscript.




```
geocode-csv.R
1 # Geocoding a csv column of "addresses" in R
2
3 #load ggmap
4 library(ggmap)
5
6 # Select the file from the file chooser
7 fileToLoad <- file.choose(new = TRUE)
8
9 # Read in the CSV data and store it in a variable
10 origAddress <- read.csv(fileToLoad, stringsAsFactors = FALSE)
11
12 # Initialize the data frame
13 geocoded <- data.frame(stringsAsFactors = FALSE)
14
15 # Loop through the addresses to get the latitude and longitude of each address
16 # origAddress data frame in new columns lat and lon
17 for(i in 1:nrow(origAddress))
18 {
19   # Print("Working...")
20   result <- geocode(origAddress$addresses[i], output = "latlon", source = "ggmap")
21   origAddress$lon[i] <- as.numeric(result[1])
22   origAddress$lat[i] <- as.numeric(result[2])
23   origAddress$geoAddress[i] <- as.character(result[3])
24 }
25 # Write a CSV file containing origAddress to the working directory
26 write.csv(origAddress, "geocoded.csv", row.names=FALSE)
27
```

20:36 (Top Level) ↕

Comments



Comments are like helping text in your R program



These statements are ignored by the interpreter while executing your actual program.



Single comment is written using # in the beginning of the statement.

#It is considered good code to have your scripts commented.

Setting Variables

- When you define a variable at the command prompt, the variable is held in your *workspace*.
- The workspace is held in the computer's main memory but can be saved to disk when you exit from R.
- The variable definition remains in the workspace until removed.
- R is a *dynamically typed language*, which means that we can change a variable's data type at will.
- We could set a to be numeric, and then turn around and immediately overwrite that with (by example) a vector of character strings. eg;
 - `> a <- 4`
 - `> a <- c("sun", "moon", "rain", "clouds")`

Printing using the print() function

- One can use the print() function which displays same result;
 - `> print(pi)`
 - `[1] 3.141593`
 - `> print(sqrt(2))`
 - `[1] 1.414214`
 - `> print(x)`
 - `[1] 3`
- The print function has a significant limitation, it prints only one object at a time.
- The only way to print multiple items is to print them one at a time.
 - `> print("The zero occurs at"); print(2*pi); print("radians")`
 - `[1] "The zero occurs at"`
 - `[1] 6.283185`
 - `[1] "radians"`

Printing using the cat() function

- The cat() function is an alternative to print that lets you combine multiple items;
 - `> cat("The zero occurs at", 2*pi, "radians.", "\n")`
 - The zero occurs at 6.283185 radians.
- By default, the cat() function puts a space between each item.

Shortcut Keys

Labeled key	Ctrl-key combination	Effect
Up arrow	Ctrl-P	Recall previous command by moving backward through the history of commands
Down arrow	Ctrl-N	Move forward through the history of commands
Backspace	Ctrl-H	Delete the character to the left of cursor
Delete	Ctrl-D	Delete the character to the right of cursor
Home	Ctrl-A	Move cursor to the start of the line
End	Ctrl-E	Move cursor to the end the line
Rsish arrow	Ctrl-F	Move cursor right (forward) one character
Left arrow	Ctrl-B	Move cursor left (back) one character
	Ctrl-K	Delete everything from the cursor position to the end of the line.
	Ctrl-U	Clear the whole line and start over
Tab		Name completion (used on some platforms)

Getting help on a function

```
> help(functionname)
```

```
> args(functionname)
```

```
> example(functionname)
```


Getting Help in R

- The command `help.start()` or choosing HTML Help from the Help menu will yield a table of contents that points to help files, manuals, frequently asked questions and the like.
- To get help for a certain function such as `summary`, use `help(summary)` or prefix the topic with a question mark: `?summary`.
- To get help on an operator, enclose it in quotes as in `help("<-")` for the assignment operator.
- If you don't know the name of a command or operator, use `help.search("your search string")` to search the built-in help files.
- Search the R site using `RSiteSearch("your search string")` or go to <http://www.r-project.org/> and click search.

The Working Directory

- The working directory is the default location for all input and output
 - Reading and writing data files
 - Opening and saving script files
 - Saving workspace image
- From the command line,
 - `getwd()` – reports the working directory
 - `setwd()` – changes the working directory
 - `setwd("C:/myfolder/data")`
 - `setwd("C :\\myfolder\\data")`



Downloading files

- From the internet:
 - `download.file("http://www..xls", "C:/test/students.xls", method="auto", quiet=FALSE, mode = "wb", cacheOK = TRUE)`

Search Path – list of packages currently loaded into the memory

```
> search( )
```

```
[1] ".GlobalEnv"      "tools:rstudio"    "package:stats"
```

```
[4] "package:graphics" "package:grDevices" "package:utils"
```

```
[7] "package:datasets" "package:methods"  "Autoloads"
```

```
[10] "package:base"
```



Variables

- Variables store values and are an important component in programming. A variable can store a number, an object, a statistical result, vector, dataset, a model prediction basically anything R outputs
- A valid variable name consists of letters, numbers and the dot or underline characters.
- The variable name starts with a letter or the dot should not follow by a number.

Variable Name	Validity	Reason
var_name2.	valid	Has letters, numbers, dot and underscore
var_name%	Invalid	Has the character '%'. Only dot(.) and underscore allowed.
2var_name	invalid	Starts with a number
.var_name, var.name	valid	Can start with a dot(.) but the dot(.)should not be followed by a number.
.2var_name	invalid	The starting dot is followed by a number making it invalid.
_var_name	invalid	Starts with _ which is not valid

Variable Assignment

- To declare a variable, we need to assign a variable name. The name should not have space. We can use _ to connect to words.
- To add a value to the variable, use <- or =.


Here is the syntax:

```
# First way to declare a variable: use the '<-'  
name_of_variable <- value
```


```
# Second way to declare a variable: use the '='  
name_of_variable = value
```



Terms

- A variable is a quantity, quality, or property that you can measure.
- A value is the state of a variable when you measure it. The value of a variable may change from measurement to measurement.
- An observation is a set of measurements made under similar conditions (you usually make all of the measurements in an observation at the same time and on the same object). An observation will contain several values, each associated with a different variable. I'll sometimes refer to an observation as a data point.
- Tabular data is a set of values, each associated with a variable and an observation. Tabular data is tidy if each value is placed in its own “cell”, each variable in its own column, and each observation in its own row.



Data Types in R





Types

Converting between common data types in R. Can always go from a higher value in the table to a lower value.

as.logical	TRUE, FALSE, TRUE	Boolean values (TRUE or FALSE).
as.numeric	1, 0, 1	Integers or floating point numbers.
as.character	'1', '0', '1'	Character strings. Generally preferred to factors.
as.factor	'1', '0', '1', Levels: '1', '0'	Character strings with preset levels. Needed for some statistical models.

Maths Functions

log(x)	Natural log.	sum(x)	Sum.
exp(x)	Exponential.	mean(x)	Mean.
max(x)	Largest element.	median(x)	Median.
min(x)	Smallest element.	quantile(x)	Percentage quantiles.
round(x, n)	Round to n decimal places.	rank(x)	Rank of elements.
signif(x, n)	Round to n significant figures.	var(x)	The variance.
cor(x, y)	Correlation.	sd(x)	The standard deviation.

Variable Assignment

```
> a <- 'apple'
> a
[1] 'apple'
```







The Environment

ls()	List all variables in the environment.
rm(x)	Remove x from the environment.
rm(list = ls())	Remove all variables from the environment.

You can use the environment panel in RStudio to browse variables in your environment.




Matrices

```
m <- matrix(x, nrow = 3, ncol = 3)
# Create a matrix from x.
```

 m[2,]	- Select a row	 t(m)	Transpose
 m[, 1]	- Select a column	 m %*% n	Matrix Multiplication
 m[2, 3]	- Select an element	 solve(m, n)	Find x in: m * x = n

Lists

```
l <- list(x = 1:5, y = c('a', 'b'))
# A list is a collection of elements which can be of different types.
```

 l[[2]]	 l[[1]]	 l\$x	 l['y']
Second element of l.	New list with only the first element.	Element named x.	New list with only element named y.




Also see the [dplyr](#) package.

Data Frames



```
df <- data.frame(x = 1:3, y = c('a', 'b', 'c'))
# A special case of a list where all elements are the same length.
```

x	y
1	a
2	b
3	c

Matrix subsetting

df[, 2]	
df[2,]	
df[2, 2]	

List subsetting

df\$x		df[[2]]	
Understanding a data frame			
View(df)	See the full data frame.		
head(df)	See the first 6 rows.		

nrow(df)
Number of rows.


ncol(df)
Number of columns.

dim(df)
Number of columns and rows.

cbind - Bind columns.



rbind - Bind rows.



Strings

Also see the [stringr](#) package.

paste(x, y, sep = ' ')	Join multiple vectors together.
paste(x, collapse = ' ')	Join elements of a vector together.
grep(pattern, x)	Find regular expression matches in x.
gsub(pattern, replace, x)	Replace matches in x with a string.
toupper(x)	Convert to uppercase.
tolower(x)	Convert to lowercase.
nchar(x)	Number of characters in a string.

Factors

factor(x)	Turn a vector into a factor. Can set the levels of the factor and the order.
cut(x, breaks = 4)	Turn a numeric vector into a factor by 'cutting' into sections.

Statistics

lm(y ~ x, data=df)	Linear model.	t.test(x, y)	Perform a t-test for difference between means.	prop.test	Test for a difference between proportions.
glm(y ~ x, data=df)	Generalised linear model.	pairwise.t.test	Perform a t-test for paired data.	aov	Analysis of variance.
summary	Get more detailed information out a model.				

Distributions

	Random Variates	Density Function	Cumulative Distribution	Quantile
Normal	rnorm	dnorm	pnorm	qnorm
Poisson	rpois	dpois	ppois	qpois
Binomial	rbinom	dbinom	pbinom	qbinom
Uniform	runif	dunif	punif	qunif

Plotting

Also see the [ggplot2](#) package.

 plot(x)	Values of x in order.	 plot(x, y)	Values of x against y.	 hist(x)	Histogram of x.
---	-----------------------	--	------------------------	---	-----------------

Dates

See the [lubridate](#) package.

Declare variables of different types

Example 1:

```
# Numeric
```

```
x <- 28
```

```
class(x)
```

```
Output: ## [1] "numeric"
```

Example 2:

```
# String
```

```
y <- "R is Fantastic"
```

```
class(y)
```

```
Output: ## [1] "character"
```

Example 3:

```
# Boolean
```

```
z <- TRUE
```

```
class(z)
```

```
Output: ## [1] "logical"
```


In the command line, we can write the following codes to see what happens:

Example 1:

```
# Print variable x
```

```
x <- 42
```

```
x
```

```
Output: ## [1] 42
```

Example 2:

```
y <- 10
```

```
y
```

```
Output: ## [1] 10
```

Example 3:

```
# We call x and y and apply a subtraction
```

```
x-y
```

```
Output: ## [1] 32
```

Basic Arithmetic Operators

Operator	Description
+	Addition
-	Subtraction
*	Multiplication
/	Division
^ or **	Exponentiation

Examples:

A multiplication

$3*9$

Output: ## [1] 27

A division

$(5+13)/2$

Output: ## [1] 9

Exponentiation

-3^2

Output: ## [1] 9

A Function

- Create a function by using the function keyword followed by a list of parameters and the function body. A one-liner looks like this:
 - `function(param1, ..., paramN) expr`
- The function body can be a series of expressions, in which case curly braces should be used around the function body:
 - `function(param1, ..., paramN) {`
 - `expr1`
 - `.`
 - `.`
 - `.`
 - `exprM`
 - `}`

Defining a Function

- Create a function by using the function keyword followed by a list of parameters and the function body. A one-liner looks like this:
 - `function(param1, ..., paramN) expr`
- The function body can be a series of expressions, in which case curly braces should be used around the function body:
 - `function(param1, ..., paramN) {`
 - `expr1`
 - `.`
 - `.`
 - `.`
 - `exprM`
 - `}`

Defining a Function

- e.g. function for calculating the coefficient of variation.

```
> cv <- function(x) sd(x)/mean(x)  
> cv(1:18)  
[1] 0.5619515
```

- The first line creates a function and assigns it to cv.
- The second line invokes the function, using 1:18 for the value of parameter x.

Logical

- These are applicable only to vectors of type logical, numeric or complex.
- All numbers greater than 1 are considered as logical value TRUE.
- Each element of the first vector is compared with the corresponding element of the second vector.
- The result of comparison is a Boolean value.

Data Structures in



VECTORS

1

MATRIX

2

ARRAY

3

LIST

4

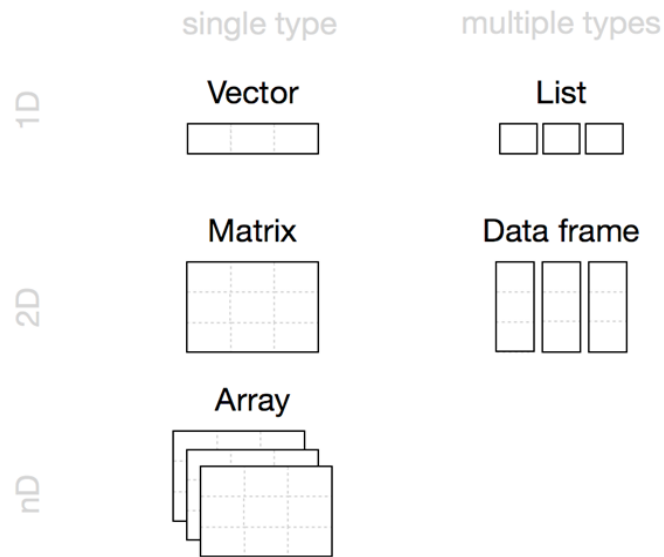
DATA FRAME

5

Learning objectives of this module:

- What are the three properties of a vector, other than its contents?
- What are the four common types of atomic vectors? What are the two rare types?
- What are attributes? How do you get them and set them?
- How is a list different from an atomic vector? How is a matrix different from a data frame?
- Can you have a list that is a matrix?

Fundamental Data Structures



Vector

- A sequence of numbers or characters, or higher-dimensional arrays like matrices

Matrix

- The basic two dimensional data structure in R is the vector

Array

- If a higher dimension vector is desired, then use the function to generate the n-dimensional object.

List


- An ordered set of components stored in a 1D vector

Data.Frame

- A table-like structure (experimental results often collected in this form)

Factor

- A sequence assigning a category to each index.



Atomic Data Elements: Vectors

- In R the “base” type is a vector, not a scalar.
- A vector is an indexed set of values that are all of the same type. The type of the entries determines the class of the vector.
- The possible vectors data types are:
 - integer
 - numeric
 - character
 - complex
 - logical
- An integer is a subclass of numeric.
- Cannot combine vectors of different modes

Creating Vectors

Assignment of a value to a variable is done with `<-` (two symbols, no space).

```
> v <- 1
```

```
➤ v
```

Vectors can only contain entries of the same type: numeric or character; you can't mix them.

- Note that characters should be surrounded by `" "`. The most basic way to create a vector is with `c(x1, . . . , xn)`, and it works for characters and numbers alike.

```
> x <- c("a", "b", "c")
```

```
> length(x)
```

```
[1] 3
```

Can also generate regular sequences:

seq(from = #, to = #, by = #): allows you to create a sequence from a starting number to an ending number.

rep(x = c(1, 2), each = 3): function allows you to repeat a scalar (or vector) a specified number of times, or to a desired length.

Selecting Vector Elements

You can select the indexing technique appropriate:

- Use square brackets [] to select vector elements by their position
v[2] – second element of v
- Use negative indexes to exclude elements
- Use a vector of indexes to select multiple values.
- Use a logical vector to select elements based on a condition
- Use names to access named elements.

Vector Arithmetic

Numeric vectors can be used in arithmetic expressions, in which case the operations are performed element by element to produce another vector.

```
> x <- rnorm(3)
```

```
> y <- rnorm(3)
```

```
> x
```

```
> x + 1
```

```
> c(1, 2, 3, 4, 5) + c(5, 4, 3, 2, 1)
```


More Vector Arithmetic Statistical operations on numeric vectors

- In studying data, you will make frequent use of sum, which gives the sum of the entries, max, min, mean.

Function	Example	Result
sum(x), product(x)	sum(1:20)	210
min(x), max(x)	min(1:20)	1
mean(x), median(x)	mean(1:20)	10.5
sd(x), var(x), range(x)	sd(1:20)	5.91608
quantile(x, probs)	quantile(1:20, probs = .2)	20%, 4.8
summary(x)	summary(1:20)	Min = 1.00. 1st Qu. = 5.75, Median = 10.50, Mean = 10.50, 3rd Qu. = 15.25, Max = 20.0

> (i.e., $\text{Sum}((x - \text{mean}(x))^2)/(\text{length}(x) - 1)$)

- A useful function for quickly getting properties of a vector: summary(y)

More Vector Arithmetic Statistical operations on continuous vectors

Function	Description	Example	Result
<code>round(x, digits)</code>	Round elements in x to digits digits	<code>round(c(3.712, 3.1415), digits = 1)</code>	3.7, 3.1
<code>ceiling(x), floor(x)</code>	Round elements x to the next highest (or lowest) integer	<code>ceiling(c(2.6, 8.1))</code>	3, 9
<code>x %% y</code>	Modular arithmetic (ie. $x \bmod y$)	<code>8 %% 4</code>	0

Vector Arithmetic Statistical operations on discrete vectors

Function	Description	Example	Result						
unique(x)	Returns a vector of all unique values.	unique(c(3, 3, 4,5,12))	3, 4, 5, 12						
table(x, exclude)	Returns a table showing all the unique values as well as a count of each occurrence. To include a count of NA values, include the argument exclude = NULL	table(c("x", "x", "y", "z"))	<table><tr><td>x</td><td>y</td><td>z</td></tr><tr><td>2</td><td>1</td><td>1</td></tr></table>	x	y	z	2	1	1
x	y	z							
2	1	1							

#See the “indexing vectors.R” in the folder.

#Create the dataframe using the details in the below table

baby.names	baby.city	baby.ages	baby.weight	baby.eyecolor
amy	macon	13	21	brown
brittany	athens	21	22	brown
carol	pink	32	41	green
donna	savannah	6	16	blue
erin	savannah	12	18	blue
fran	atlanta	11	19.4	grey
gigi	atlanta	18	26	brown
helen	athens	16	23	green
irene	macon	17	22	brown
jackie	macon	34	36	brown

#Exercise continued:
#After creating the dataframe
#Slice the dataframe to answer the below questions.

- # What was the weight of the first baby?
- # What were the ages of the first 5 babies?
- # What were the names of the babies born with green eyes?
- # What were the weights of either blue or grey eyed babies?
- # Change the age of baby "irene" to 18
- # How many babies born in canton are in the data?
- # What percent of babies are older than 14 months?

Relational Operators return values inside the vector based on logical conditions. Each element of the first vector is compared with the corresponding element of the second vector. The result of comparison is a Boolean value

Operator	Description
<	Less than
<=	Less than or equal to
>	Greater than
>=	Greater than or equal to
==	Exactly equal to
!=	Not equal to
!x	Not x
x	y
x & y	x AND y
isTRUE(x)	Test if X is TRUE

You can add many conditional statements, but we need to include them in a parenthesis. Follow this structure to create a conditional statement:

```
variable_name[(conditional_statement)]
```

Logical Operators

- These are applicable only to vectors of type logical, numeric or complex.
- All numbers greater than 1 are considered as logical value TRUE.
- Each element of the first vector is compared with the corresponding element of the second vector.
- The result of comparison is a Boolean value.

&	It is called Element-wise Logical AND operator. It combines each element of the first vector with the corresponding element of the second vector and gives a output TRUE if both the elements are TRUE.
	It is called Element-wise Logical OR operator. It combines each element of the first vector with the corresponding element of the second vector and gives a output TRUE if nay one the elements are TRUE.
!	It is called Logical NOT operator. Takes each element of the vector and gives the opposite logical value.

- The logical operator && and || considers only the first element of the vectors and give a vector of single element as output.

Example:

Create a vector from 1 to 8

```
logical_vector <- c(1:8)
```

```
logical_vector > 6
```

Output: ## [1] FALSE FALSE FALSE FALSE FALSE FALSE TRUE TRUE

Example:

```
logical_vector <- c(1:8)
```

```
logical_vector[(logical_vector>3) & (logical_vector<5)]
```

Output: ## [1] 4

Factors in R: Categorical & Continuous Variables

Factors are variables in R which take on a limited number of different values; such variables are often referred to as categorical variables.

In a dataset, we can distinguish two types of variables:

categorical and **continuous**

- In a categorical variable, the value is limited and usually based on a particular finite group. For example, a categorical variable can be countries, year, gender, occupation.
- A continuous variable, however, can take any values, from integer to decimal. For example, we can have the revenue, price of a share, etc..

Categorical Variables

R stores categorical variables into a factor. Let's check the code below to convert a character variable into a factor variable. Characters are not supported in a machine learning algorithm, and the only way is to convert a string to an integer.

Syntax

```
factor(x = character(), levels, labels = levels, ordered = is.ordered(x))
```

Arguments:

- **x**: A vector of data. Need to be a string or integer, not decimal.
- **Levels**: A vector of possible values taken by x. This argument is optional. The default value is the unique list of items of the vector x.
- **Labels**: Add a label to the x data. For example, 1 can take the label `male` while 0, the label `female`.
- **ordered**: Determine if the levels should be ordered.

Nominal Categorical Variable

A categorical variable has several values but the order does not matter. For instance, male or female categorical variable do not have ordering.

```
# Create a color vector
```

```
color_vector <- c('turquoise', 'red', 'green', 'ivory', 'black', 'yellow')
```

```
# Convert the vector to factor
```

```
factor_color <- factor(color_vector)
```

```
factor_color
```

Output:

```
## [1] turquoise red green ivory black yellow
```

```
## Levels: black turquoise green red ivory yellow
```

Ordinal Categorical Variable

Ordinal categorical variables do have a natural ordering. We can specify the order, from the lowest to the highest with `order = TRUE` and highest to lowest with `order = FALSE`.

Example: We can use `summary` to count the values for each factor.

```
# Create Ordinal categorical vector
day_vector <- c('evening', 'morning', 'afternoon', 'midday', 'midnight', 'evening')
# Convert `day_vector` to a factor with ordered level
factor_day <- factor(day_vector, order = TRUE, levels = c('morning', 'midday', 'afternoon', 'evening', 'midnight'))
# Print the new variable
factor_day
```

Output:

```
## [1] evening morning afternoon midday midnight evening
```

Example:

```
## Levels: morning < midday < afternoon < evening < midnight
# Append the line to above code
# Count the number of occurrence of each level summary(factor_day)
```

Output:

```
## morning midday afternoon evening midnight
## 1 1 1 2 1
```

Lists in R:

- A list is a generic object consisting of an ordered collection of objects.
- Lists are heterogeneous data structures.
- A list is an one-dimensional data structures.
- A list can be a list of vectors, list of matrices, a list of characters and a list of functions and so on.
- Lists are different from atomic vectors because their elements can be of any type, including lists.
- You construct lists by using `list()` instead of `c()`:

```
x <- list(1:5, "d", c(FALSE, FALSE, TRUE), c(3.5 7.2))  
str(x)
```

Example - Illustrate a List

The first attributes is a numeric vector
containing the employee IDs which is
created using the 'c' command here
empld = c(1, 2, 3, 4)

The second attribute is the employee's name
which is created using this line of code here
which is the character vector
empName = c("Ann", "Sanjan", "Ellison", "Evette")

The third attribute is the number of employees
which is a single numeric variable.
numberOfEmp = 4

We can combine all these three different
data types into a list
containing the details of employees
which can be done using a list command
empList = list(empld, empName, numberOfEmp)

print(empList)

Matrix – What is a Matrix?

- A matrix is a 2-dimensional array that has m number of rows and n number of columns. In other words, matrix is a combination of two or more vectors with the same data type.

$\begin{bmatrix} 1 & 5 \\ -3 & 6 \end{bmatrix}$	$\begin{bmatrix} -1 & 5 \\ 4 & 7 \\ -8 & 2 \end{bmatrix}$	$\begin{bmatrix} 3 \\ 10 \\ -1 \end{bmatrix}$	$[-2 \ 4 \ 7 \ -6]$
(2 x 2)	(3 x 2)	(3 x 1)	(1 x 4)

- You can create a matrix with the function `matrix()`. This function takes three arguments:
`matrix(data, nrow, ncol, byrow = FALSE)`

Arguments:

- `data`: The collection of elements that R will arrange into the rows and columns of the matrix.
- `nrow`: Number of rows
- `ncol`: Number of columns
- `byrow`: The rows are filled from the left to the right. We use ``byrow = FALSE`` (default values), if we want the matrix to be filled by the columns i.e. the values are filled top to bottom.

#Print dimension of the matrix with dim()

Defining names of columns and rows in a matrix

```
# Print dimension of the matrix with dim()  
dim(matrix_a)
```

In order to define rows and column names, you can create two vectors of different names, one for row and other for a column. Then, using the Dimnames attribute, you can name them appropriately:

```
rows = c("row1", "row2", "row3", "row4")    #Creating our character vector of row names
```

```
cols = c("coln1", "coln2", "coln3")          #Creating our character vector of column names
```

```
mat <- matrix(c(4:15), nrow = 4, byrow = TRUE, dimnames = list(rows, cols) )
```

```
#creating our matrix mat and assigning our vectors to dimnames
```

```
# Print matrix  
print(mat)
```


Add a Column to a Matrix with the cbind()

- You can add a column to a matrix with the cbind() command.
- cbind() means column binding. cbind() can concatenate as many matrix or columns as specified.

Example:

```
# concatenate c(1:5) to the matrix_a  
matrix_a1 <- cbind(matrix_a, c(1:5))  
# Check the dimension dim(matrix_a1)
```

Output:

```
## [1] 5 3
```

Example:

```
matrix_a1
```

Output

```
##      [,1] [,2] [,3]  
## [1,]    1    2    1  
## [2,]    3    4    2  
## [3,]    5    6    3  
## [4,]    7    8    4  
## [5,]    9   10    5
```

Example:

We can also add more than one column. Add the next sequence of number to the matrix_a2 matrix. The dimension of the new matrix will be 4x6 with number from 1 to 24.

```
matrix_a2 <- matrix(13:24, byrow = FALSE, ncol = 3)
```

```
##      [,1] [,2] [,3]  
## [1,]  13  17  21  
## [2,]  14  18  22  
## [3,]  15  19  23  
## [4,]  16  20  24
```

Example:

```
matrix_c <- matrix(1:12, byrow = FALSE, ncol = 3)  
matrix_d <- cbind(matrix_a2, matrix_c)  
dim(matrix_d)
```

Output:

```
## [1] 4 6
```

Slice a Matrix: Accessing individual components

We can select elements one or many elements from a matrix by using the square brackets `[]`. This is where slicing comes into the picture.


For example:

`matrix_c[1,2]` selects the element at the first row and second column.

`matrix_c[1:3,2:3]` results in a matrix with the data on the rows 1, 2, 3 and columns 2, 3,

`matrix_c[,1]` selects all elements of the first column.

`matrix_c[1,]` selects all elements of the first row.



Examples:

```
> A = matrix( + c(2, 4, 3, 1, 5, 7),      # the data elements
              + nrow=2,                  # number of rows
              + ncol=3,                  # number of columns
              + byrow = TRUE)           # fill matrix by rows

> A


> A[2, 3]                               # element at 2nd row, 3rd column
> A[2, ]                                 # the 2nd row
> A[,c(1,3)]                            # the 1st and 3rd columns
> t(A)                                  # transpose of B
> c(B)                                  # deconstruct a matrix
```

Functions for viewing matrices and dataframes and returning information about them.

Function	Description
head(x), tail(x)	Print the first few rows (or last few rows).
View(x)	Open the entire object in a new window
nrow(x), ncol(x), dim(x)	Count the number of rows and columns
rownames(), colnames(), names()	Show the row (or column) names
str(x), summary(x)	Show the structure of the dataframe (ie., dimensions and classes) and summary statistics

Introduction to Arrays in R

- In arrays, data is stored in the form of matrices, rows, and columns.



R Array Syntax

`Array_NAME <- array(data, dim = (row_Size, column_Size, matrices, dimnames)`

- **data** – Data is an input vector that is given to the array.
- **matrices** – Array in R consists of multi-dimensional matrices.
- **row_Size** – row_Size describes the number of row elements that an array can store.
- **column_Size** – Number of column elements that can be stored in an array.
- **dimnames** – Used to change the default names of rows and columns to the user's preference.

Arguments in Array

The array function in R can be written as:

`array(data = NA, dim = length(data), dimname = NULL)`

- **data** is a vector that provides data to fill the array.
- **dim** attribute provides maximum indices in each dimension
- **dimname** can be either NULL or can have a name for the array.

How to Create an Array in R

Example:

Create two vectors of different lengths.

```
vector1 <- c(2,8,3)
```

```
vector2 <- c(10,14,17,12,11,15)
```

Take these vectors as input to the array.

```
result <- array(c(vector1,vector2), dim = c(3,3,2))
```

```
print(result)
```


Different Operations on Rows and Columns

Naming Columns And Rows

```
# Example:
# Create two vectors of different lengths.
vector1 <- c(2,8,3)
vector2 <- c(10,14,17,12,11,15)
column.names <- c("COL1","COL2","COL3")
row.names <- c("ROW1","ROW2","ROW3")
matrix.names <- c("Matrix1","Matrix2")

# Take these vectors as input to the array.
result <- array(c(vector1,vector2),dim = c(3,3,2), dimnames = list(row.names,
column.names, matrix.names))

print(result)
```

Print the third row of the first matrix of the array.

```
print(result[3,,1])
```

Print the element in the 2nd row and 3rd column of the 2nd matrix.

```
print(result[2,3,2])
```

Print the 2nd Matrix.

```
print(result[:,,2])
```

Manipulating R Array Elements

Example:

Create two vectors of different lengths.

```
vector1 <- c(1,4,6)
```

```
vector2 <- c(2,3,5,6,7,9)
```

Take these vectors as input to the array.

```
array1 <- array(c(vector1,vector2), dim = c(3,3,2))
```

Create two vectors of different lengths.

```
vector3 <- c(6,4,1)
```

```
vector4 <- c(9,7,6,5,3,2)
```

```
array2 <- array(c(vector1,vector2), dim = c(3,3,2))
```

```
# create matrices from these arrays.
```

```
matrix1 <- array1[,2]
```

```
matrix2 <- array2[,2]
```

```
# Add the matrices.
```

```
result <- matrix1+matrix2
```

```
print(result)
```

Calculations across R Array Elements

- `apply()` function for calculations in an array in R.
- Syntax `apply(x, margin, fun)`

Following is the description of the parameters used:

- `x` is an array.
- `margin` is the name of the dataset used.
- `fun` is the function to be applied to the elements of the array.
- For Example:
 - We use the `apply()` function below in different ways. To calculate the sum of the elements in the rows of an array across all the matrices.

Example:

We will create two vectors of different lengths.

```
vector1 <- c(1,4,6)
```

```
vector2 <- c(2,3,5,6,7,9)
```

Now, we will take these vectors as input to the array.

```
new.array <- array(c(vector1,vector2),dim = c(3,3,2))
```

```
print(new.array)
```

Use (apply) to calculate the sum of the rows across all the matrices.

```
result <- apply(new.array, c(1), sum)
```

```
print(result)
```

R Data Frames:

- Data frames combine the behavior of lists and matrices to make a structure ideally suited for the needs of statistical data. The data frame is a list of vectors which are of equal length.
- A data-frame must have column names and every row should have a unique name.
 - `names()`, `colnames()`, and `rownames()`
- Each column must have the identical number of items.
- Each item in a single column must be of the same data type.
- Different columns may have different data types.
- A matrix contains only one type of data, while a data frame accepts different data types (numeric, character, factor, etc.). This makes it a 2-dimensional structure, so it shares properties of both the matrix and the list.

How to Create a Data Frame

- We can create a data frame by passing the variable a,b,c,d into the `data.frame()` function. We can name the columns with `name()` and simply specify the name of the variables.
- `data.frame(df, stringsAsFactors = TRUE)`

Arguments:

- `df`: It can be a matrix to convert as a data frame or a collection of variables to join
- `stringsAsFactors`: Convert string to factor by default


```
df <- data.frame(x = 1:3, y = c("a", "b", "c"))  
str(df)
```

```
#> 'data.frame': 3 obs. of 2 variables:  
#> $ x: int 1 2 3  
#> $ y: Factor w/ 3 levels "a","b","c": 1 2 3
```

```
df <- data.frame(  
  x = 1:3,  
  y = c("a", "b", "c"),  
  stringsAsFactors = FALSE)  
str(df)
```

```
#> 'data.frame': 3 obs. of 2 variables:  
#> $ x: int 1 2 3  
#> $ y: chr "a" "b" "c"
```

Because a data.frame is an S3 class, its type reflects the underlying vector used to build it: the list. To check if an object is a data frame, use `class()` or test explicitly with `is.data.frame()`:

```
typeof(df)  
#> [1] "list"  
class(df)  
#> [1] "data.frame"  
is.data.frame(df)  
#> [1] TRUE
```

You can coerce an object to a data frame with `as.data.frame()`:

- A vector will create a one-column data frame.
- A list will create one column for each element; it's an error if they're not all the same length.
- A matrix will create a data frame with the same number of columns and rows as the matrix.

```
# Example R program to illustrate dataframe

# A vector which is a character vector
Name = c("Auriel", "Ray", "Asia")

# A vector which is a character vector
Type = c("O+", "B-", "A-")

# A vector which is a numeric vector
Age = c(36, 23, 62)

# To create dataframe use data.frame command
# and then pass each of the vectors
# we have created as arguments
# to the function data.frame()

df = data.frame(Name, Type, Age)

print(df)
```

Adding new columns through simple list-like assignments.

```
> x$State <- c("KY", "PA")
> x
  SN Age Name County State
1  24 Felicia  Jefferson KY
2  16 Mary    Montreal  PA
```

Data frame columns can be **deleted** by assigning NULL to it.

```
> x$State <- NULL
> x
  SN Age Name County
1  24 Felicia  Jefferson
2  16 Mary    Montreal
```

Similarly, rows can be **deleted** through assignments.

```
> x <- x[-1,]
> x
  SN Age Name County
1  24 Felicia  Jefferson
2  16 Mary    Montreal
```

How to access components of a data frame?

Assessing like a list:

We can use either `[`, `[[` or `$` operator to access columns of data frame.

```
x["Name"]
Name
1 Auriel
2 Asia
> x$Name
[1] "Auriel" "Asia"
> x[["Name"]]
[1] "Auriel" "Asia"
> x[[3]]
[1] "Auriel" "Asia"
```

How to access components of a data frame?

Assessing like a matrix:

Data frames can be accessed like a matrix by providing index for row and column.

We can use either `[], [[` or `$` operator to access columns of data frame.

```
x["Name"]
Name
1 Auriel
2 Asia
> x$Name
[1] "Auriel" "Asia"
> x[["Name"]]
[1] "Auriel" "Asia"
> x[[3]]
[1] "Auriel" "Asia"
```

It's a common mistake to try and create a data frame by `cbind()`ing vectors together. This doesn't work because `cbind()` will create a matrix unless one of the arguments is already a data frame. Instead use `data.frame()` directly:

```
bad <- data.frame(cbind(a = 1:2, b = c("a", "b")))
```

```
str(bad)
```

```
#> 'data.frame': 2 obs. of 2 variables:
```

```
#> $ a: Factor w/ 2 levels "1","2": 1 2
```

```
#> $ b: Factor w/ 2 levels "a","b": 1 2
```

```
good <- data.frame(a = 1:2, b = c("a", "b"), stringsAsFactors = FALSE)
```

```
str(good)
```

```
#> 'data.frame': 2 obs. of 2 variables:
```

```
#> $ a: int 1 2
```

```
#> $ b: chr "a" "b"
```

Merging: Combine data frames using cbind() and rbind()

```
cbind(df, data.frame(z = 3:1))
```

```
#>   x y z
```

```
#> 1 1 a 3
```

```
#> 2 2 b 2
```

```
#> 3 3 c 1
```

```
rbind(df, data.frame(x = 10, y = "z"))
```

```
#>   x y
```

```
#> 1 1 a
```

```
#> 2 2 b
```

```
#> 3 3 c
```

```
#> 4 10 z
```

- When combining column-wise, the number of rows must match, but row names are ignored.
- When combining row-wise, both the number and names of columns must match.
- Use `plyr::rbind.fill()` to combine data frames that don't have the same columns.

Example:

Let's create a factor data frame.

```
# Create gender vector
```

```
gender_vector <- c("Male", "Female", "Female", "Male", "Male")
```

```
class(gender_vector)
```

```
# Convert gender_vector to a factor
```

```
factor_gender_vector <- factor(gender_vector)
```

```
class(factor_gender_vector)
```

Output:

```
## [1] "character"
```

```
## [1] "factor"
```

Examine a Data Frame in R with 7 Basic Functions

- [dim\(\)](#): shows the dimensions of the data frame by row and column
- [str\(\)](#): shows the structure of the data frame
- [summary\(\)](#): provides summary statistics on the columns of the data frame
- [colnames\(\)](#): shows the name of each column in the data frame
- [head\(\)](#): shows the first 6 rows of the data frame
- [tail\(\)](#): shows the last 6 rows of the data frame
- [View\(\)](#): shows a spreadsheet-like display of the entire data frame

Fixing a broken data frame

Some useful functions:

- `?read.csv`
- `head()`
- `str()`
- `class()`
- `unique()`
- `levels()`
- `which()`
- `droplevels()`

Note: For these functions you have to put the name of the data object in the parentheses (i.e., `head(CO2)`).

Also remember that you can use “?” to look up help for a function (i.e. `?str`).

Knowledge Check:

1. What are the three properties of a vector, other than its contents?

The three properties of a vector are type, length, and attributes.

2a. What are the four common types of atomic vectors?

1. Logical
2. Integer
3. Double (sometimes called numeric)
4. Character

2b. What are the two rare types?

1. Complex
2. Raw

Knowledge Check (cont.)

3a. How is a list different from an atomic vector?

The elements of a list can be any type (even a list); the elements of an atomic vector are all of the same type.

3b. How is a matrix different from a data frame?

Similarly, every element of a matrix must be the same type; in a data frame, the different columns can have different types

4. Can you have a list that is a matrix?

You can make “list-array” by assigning dimensions to a list.

Decision making

R provides the following types of decision-making statements.

Statement	Description
<i>if</i> statement	An if statement consists of a Boolean expression followed by one or more statements.
<i>if...else</i> statement	An if statement can be followed by an optional else statement, which executes when the Boolean expression is false.
switch statement	A switch statement allows a variable to be tested for equality against a list of values.

Decision making

- **The if...else if...else ladder**
 - The basic syntax

```
if(boolean_expression 1) {  
    // Executes when the boolean expression 1 is true.  
}else if( boolean_expression 2) {  
    // Executes when the boolean expression 2 is true.  
}else if( boolean_expression 3) {  
    // Executes when the boolean expression 3 is true.  
}else {  
    // executes when none of the above condition is true.  
}
```

Decision making

Switch Statement

- Select One of a List of Alternatives
- switch evaluates EXPR and accordingly chooses one of the further arguments (in ...).
- The basic syntax

switch(expression, case1, case2, case3....)

- switch works in two distinct ways depending whether the first argument (EXPR) evaluates to a character string or a number
- If the value of expression is not a character string it is coerced to integer.
- If the value of the integer is between 1 and nargs()-1 (The max number of arguments) then the corresponding element of case condition is evaluated, and the result returned.
- If expression evaluates to a character string, then that string is matched (exactly) to the names of the elements.
- If there is more than one match, the first matching element is returned.
- In the case of no match, if there is a unnamed element of ... its value is returned.

```
x <- switch( 5, "morning", "midmorning", "noon", "midday", "evening", "night" )  
print(x)
```


Loops

R programming language provides the following kinds of loop statements

Statement	Description
repeat loop	Executes a sequence of statements multiple times and abbreviates the code that manages the loop variable.
while loop	Repeats a statement or group of statements while a given condition is true. It tests the condition before executing the loop body.
for loop	Like a while statement, except that it tests the condition at the end of the loop body.

Control Statement & Description

break statement	Terminates the loop statement and transfers execution to the
the	statement immediately following the loop.

Next statement	The next statement simulates the behavior of R switch.
----------------	--

Loops

Repeat Loop

The Repeat loop executes the same code again and again until a stop condition is met.

The basic syntax

```
repeat {  
    commands  
    if(condition){  
        break  
    }  
}
```

Example;

```
v <- c("Afternoon","nap")  
cnt <- 2  
repeat{  
    print(v)  
    cnt <- cnt+1  
    if(cnt > 5){  
        break  
    }  
}
```

Loops

While Loop

The While loop executes the same code again and again until a stop condition is met.

The basic syntax

```
while (test_expression) {  
    statement  
}
```

Example;

```
v <- c("Wide","while loop")  
cnt <- 2  
while (cnt < 7){  
    print(v)  
    cnt = cnt + 1  
}
```

Loops

For Loop

A for loop is a repetition control structure that allows you to efficiently write a loop that needs to execute a specific number of times.

The basic syntax

```
for (value in vector) {  
    statements  
}
```

```
Example;  
v <- LETTERS[1:4]  
for ( i in v) {  
    print(i)  
}
```

Using packages

1

```
install.packages("readr")
```

Downloads files to computer

1 x per computer

2

```
library("readr")
```

Loads package

1 x per R Session

Vignettes

```
cointoss/  
├── .Rbuildignore  
├── cointoss.Rproj  
├── DESCRIPTION  
├── NAMESPACE  
├── R/  
├── man/  
├── tests/ Vignette files  
└── vignettes/  
    └── introduction.Rmd
```

Importing your data

R allows for the import of different data formats using specific packages that can make your job easier:

Package	Purpose
readr	Importing flat files
readxl	Getting excel files into R
haven	Import SAS, STATA and SPSS data files into R.
RMySQL and Rpostgresql	Connect to databases, access and manipulate via DBI
rvest	web scraping

Example: `mydata=read.csv(file="c:/.....csv", header=TRUE)`
`data <- read_excel('input.xlsx', sheetIndex = 1)`

Additional resources for R programming basics


<http://www.sthda.com/english/wiki/easy-r-programming-basics>


<https://www.rstudio.com/resources/webinars/data-wrangling-with-r-and-rstudio/>

<https://github.com/rstudio/webinars/blob/master/05-Data-Wrangling-with-R-and-RStudio/wrangling-webinar.pdf>



Steps of Data Analysis






Module 2:

Exploratory Data Analysis:

an approach to analyzing data sets to summarize their main characteristics, often with visual methods. - Wikipedia

Variation

- The tendency of the values of a variable to change from measurement to measurement.
- You can see variation easily in real life; if you measure any continuous variable twice, you will get two different results. This is true even if you measure quantities that are constant, like the speed of light.
- Each of your measurements will include a small amount of error that varies from measurement to measurement.
- Categorical variables can also vary if you measure across different subjects (e.g. the eye colors of different people), or different times (e.g. the energy levels of an electron at different moments).
- Every variable has its own pattern of variation, which can reveal interesting information. The best way to understand that pattern is to visualize the distribution of the variable's values.



“Get to Know” the dataset

- Doing so upfront will make the rest of the project much smoother, in 3 main ways:
 1. You’ll gain valuable hints for [Data Cleaning](#).
 2. You’ll think of ideas for [Feature Engineering](#).
 3. You’ll get a "feel" for the dataset, which will help you communicate results and deliver greater impact.
- EDA should be **quick, efficient, and decisive...** not long and drawn out!
- You see, there are infinite possible plots, charts, and tables, but you only need a **handful** to "get to know" the data well enough to work with it.



What is EDA?

An approach for data analysis that employs a variety of techniques

1. Maximize insight into a data set
2. Uncover underlying structure
3. Extract important variables
4. Detect outliers and anomalies
5. Test underlying assumptions
6. Develop parsimonious models and
7. Determine optimal factor settings

EDA is a data approach.





The EDA sequence is:

Problem => Data => Analysis=> Model=> Conclusions

As opposed for a classical approach:

Problem => Data => Model=> Analysis=> Conclusions

EDA
Techniques
are generally
graphical.



EDA is majorly performed using the following methods:

Univariate visualization – provides summary statistics for each field in the raw data set

Bivariate visualization – is performed to find the relationship between each variable in the dataset and the target variable of interest

Multivariate visualization – is performed to understand interactions between different fields in the dataset

Dimensionality reduction – helps to understand the fields in the data that account for the most variance between observations and allow for the processing of a reduced volume of data.

Here is a list of all graph types that are illustrated:

- Barplot
- Boxplot
- Density Plot
- Heatmap
- Histogram
- Line Plot
- Pairs Plot
- Polygon Plot
- QQplot
- Scatterplot
- Venn Diagram

Barplot	A barplot (or barchart; bargraph) illustrates the association between a numeric and a categorical variable. The barplot represents each category as a bar and reflects the corresponding numeric value with the bar's size.	barplot(x)
Boxplot	Displays the distribution of a numerical variable based on five summary statistics: minimum non-outlier; first quartile; median; third quartile; and maximum non-outlier. Furthermore, boxplots show the positioning of outliers and whether the data is skewed.	boxplot(x)
Density Plot	A density plot (or kernel density plot; density trace graph) shows the distribution of a numerical variable over a continuous interval. Peaks of a density plot visualize where the values of numerical variables are concentrated.	plot(density(x))
Heatmap	A heatmap (or shading matrix) visualizes individual values of a matrix with colors. More common values are typically indicated by brighter reddish colors and less common values are typically indicated by darker colors.	heatmap(cbind(x, y))
Histogram	A histogram groups continuous data into ranges and plots this data as bars. The height of each bar shows the number of observations within each range.	hist(x)

Line Plot	A line plot, visualizes values along a sequence (e.g., over time). Line plots consist of an x-axis and a y-axis. The x-axis usually displays the sequence and the y-axis the values corresponding to each point of the sequence.	<code>plot(1:length(y), y, type = "l")</code>
Pairs Plot	A pairs plot is a plot matrix, consisting of scatterplots for each variable-combination of a data frame.	<code>pairs(data.frame(x, y))</code>
Polygon Plot	A polygon plot displays a plane geometric figure (i.e., a polygon) within the plot	<code>plot(1,1 col = "white", xlab="X", ylab = "Y") polygon(x= c(0.7, 1.3, 1.3, 0.8), y=c(0.6, 1.0, 1.4, 1.3), col = "#353436")</code>
QQplot	Quantile-Quantile plot; Quantile-Quantile diagram) determines whether two data sources come from a common distribution. QQ plots draw the quantiles of the two numerical data sources against each other. If both data sources come from the same distribution, the points fall on a 45-degree angle.	<code>qqplot(x,y)</code>
Scatterplot	A scatterplot (or scatter plot; scatter graph; scatter chart; scattergram; scatter diagram) displays two numerical variables with points, whereby each point represents the value of one variable on the x-axis and the value of the other variable on the y-axis.	<code>plot(x,y)</code>
Venn Diagram	A venn diagram (or primary diagram; set diagram; logic diagram) illustrates all possible logical relations between certain data characteristics. Each characteristic is represented as a circle, whereby overlapping parts of the circles illustrate elements that have both characteristics at the same time.	<code>Install.packages("VennDiagram") library("VennDiagram") plot.new() draw.single.venn(area = 10)</code>

Any
Questions





Brian Ashford

bashford@beverasolutions.com

Yvonne Phillips

yphillips@beverasolutions.com

BeVera